

Wild Rydes : Serverless Workshop

How should my app
withstand a server failing?

How can I tell if a
server has been
compromised?

How can I increase
utilization of my servers?

Which OS should my
servers run?

How much remaining
capacity do my servers have?

When should I decide to
scale up my servers?

How should I implement dynamic
configuration changes on my servers

What size servers are
right for my budget?

How will I keep my server
OS patched?

How can I control
access from my servers?

Which packages should
be baked into my server images?

Servers

(AAHHHHHHHHH!!)

How will the application
handle server hardware failure?

How will new code be
deployed to my servers?

How many users create
too much load for my servers?

What size server is
right for my performance?


Which users should have
access to my servers?

Should I tune OS settings
to optimize my application?

When should I decide to
scale out my servers?

How many servers
should I budget for?


Owning servers means dealing with ...



Operations and management



Scaling



Provisioning and utilization



Availability and fault tolerance

What to expect from this workshop

- Serverless Architectures Overview – BRIEF!
 - Overview of
 - AWS Lambda
 - Amazon API Gateway
 - Amazon DynamoDB
 - Amazon Cognito
- Workshop Breakout – Time to build

AWS compute offerings



Amazon EC2

Resizable virtual
servers in the
cloud



Amazon ECS

Container management
service for running
Docker on EC2



AWS Lambda

Serverless compute, run
code in response to
events

Why serverless architectures?

- No servers to manage and scale
- Run at scale
- Respond quickly to events
- Only pay for compute time that you use
- Developer productivity



AWS Lambda

Benefits of using Lambda

1


No Servers to Manage



Lambda automatically runs your code without requiring you to provision or manage servers. Just write the code and upload it to Lambda.

2


Continuous Scaling



Lambda automatically scales your application by running code in response to each trigger. Your code runs in parallel and processes each trigger individually, scaling precisely with the size of the workload.

3

Subsecond Metering



With Lambda, you are charged for every 100 ms your code executes and the number of times your code is triggered. You don't pay anything when your code isn't running.

AWS Lambda – How it works



Bring your own code

- Node.js, Java, Python, C#
- Java = Any JVM based language such as Scala, Clojure, etc.
- Bring your own libraries



Simple resource model

- Select memory from 128MB to 1.5GB in 64MB steps
- CPU & Network allocated proportionately to RAM
- Reports actual usage



Flexible invocation paths


- Event or RequestResponse invoke options
- Existing integrations with various AWS services



Fine grained permissions

- Uses IAM role for Lambda execution permissions
- Uses Resource policy for AWS event sources

AWS Lambda – Use Cases




Data Processing

Execute code in response to changes in data, shifts in system state, or actions by users



Backends

Execute backend logic to handle requests for web, mobile, IoT, and 3rd party APIs



Control Systems

Customize responses and response workflows to state and data changes within AWS

Amazon API Gateway

Your feedback



Managing multiple versions and stages of an API is difficult




Monitoring third-party developers' access is time consuming



Access authorization is a challenge



Traffic spikes create an operational burden



What if I don't want servers at all?

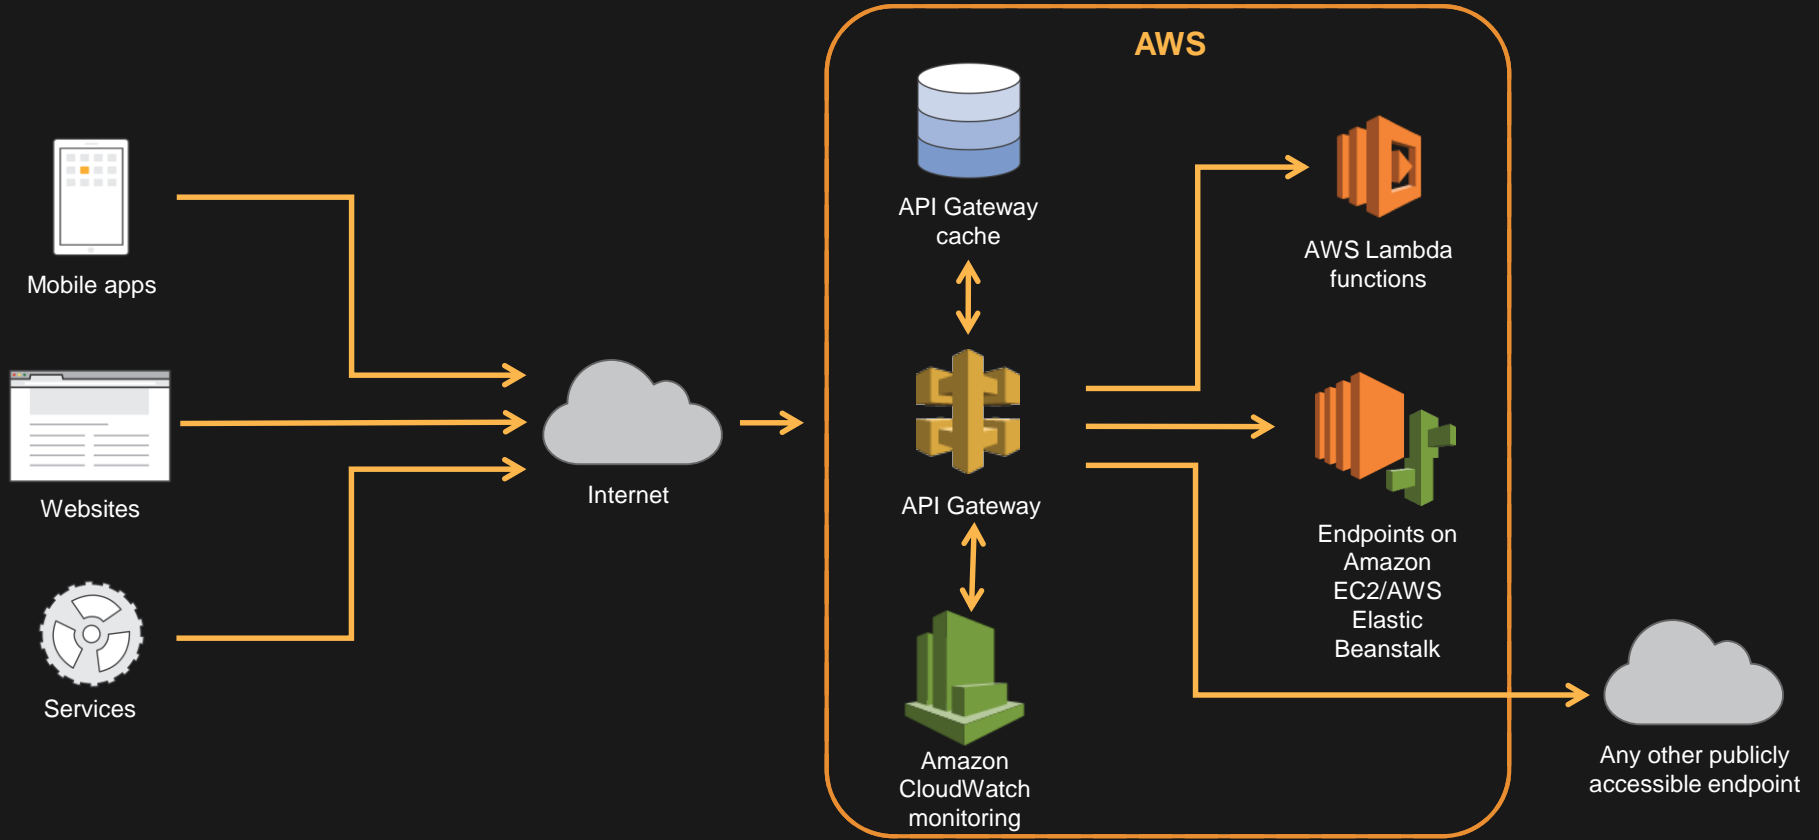
API Gateway - Capabilities

- Host multiple versions and stages of your APIs
- Create and distribute API keys to developers
- Leverage signature version 4 to authorize access to APIs
- Throttle and monitor requests to protect your backend
- Utilize Lambda as a backend

Benefits of API Gateway

- Managed cache to store API responses
- Reduced latency and distributed denial of service (DDoS) protection through Amazon CloudFront
- SDK generation for iOS, Android, and JavaScript
- Swagger support
- Request and response data transformation

An API call flow



Amazon DynamoDB

Amazon DynamoDB

Fast and flexible NoSQL database service for any scale



Dead Simple

- GetItem(primaryKey)
- PutItem(item)

```
const doc = require('dynamodb-doc');
const dynamo = new doc.DynamoDB();

exports.handler = (event, context, callback) => {

  const id = event.payload.id;
  dynamo.getItem(id, callback);

};
```

Robust Depth

- Fine-Grained Access Control
- Streams
- Triggers
- Cross-Region Replication
- DynamoDB local
- Free-text search
- Titan Graph Database integration
- Strong consistency option
- Atomic counters

Amazon Cognito

Amazon Cognito Identity



Cognito User Pools

You can easily and securely add sign-up and sign-in functionality to your mobile and web apps with a fully-managed service that scales to support 100s of millions of users.



Your own auth Guest

Federated User Identities

Your users can sign-in through social identity providers such as Facebook, Twitter and SAML providers and you can control access to AWS resources from your app.

Amazon Cognito User Pools

1

Serverless
Authentication and
User Management



Add user sign-up and sign-in easily to your mobile and web apps without worrying about server infrastructure

2

Managed User Directory



A simple, secure, low-cost, and fully managed service to create and maintain a user directory that scales to 100s of millions of users

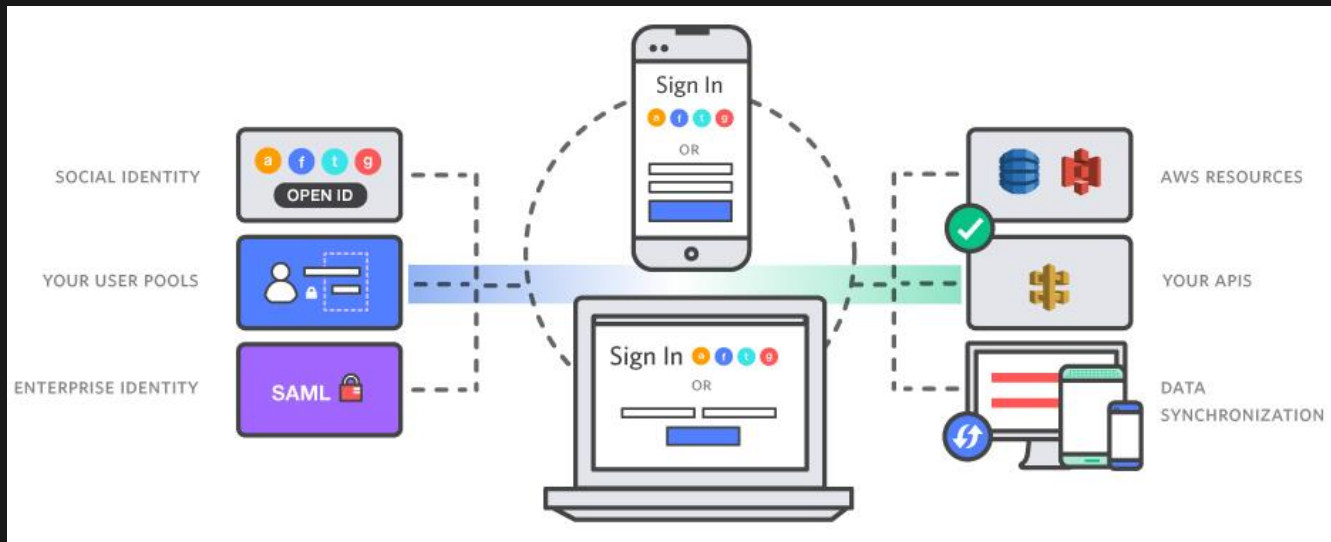
3

Enhanced Security
Features



Verify phone numbers and email addresses and offer multi-factor authentication

Comprehensive Support for Identity Use Cases



Scenario: Wild Rydes (www.wildrydes.com)



Help Wild Rydes Disrupt Transportation!

So how does this magic work?



DOWNLOAD THE APP

Head over to the app store and download the Wild Rydes app. You're just a few taps away from getting your ryde.



REQUEST A UNICORN

We can get you there. Simply request a ryde on the app and we'll connect you with a unicorn immediately.



PICK A PRICE

Pick the valuation you're willing to pay and your ryde is set up. The only surge is the acceleration you get when taking off.



RIDE OFF TO SUCCESS!

After matching with your unicorn and agreeing to its terms, you'll be all set. Your unicorn will arrive shortly to pick you up.

Wild Rydes is Backed by Leading Investors



**THE BARN
ACCELERATOR**



**TENDERLOIN
CAPITAL**



**PENGLAI COMMUNICATIONS
AND POST NEW CENTURY
TECHNOLOGY CORP LIMITED**

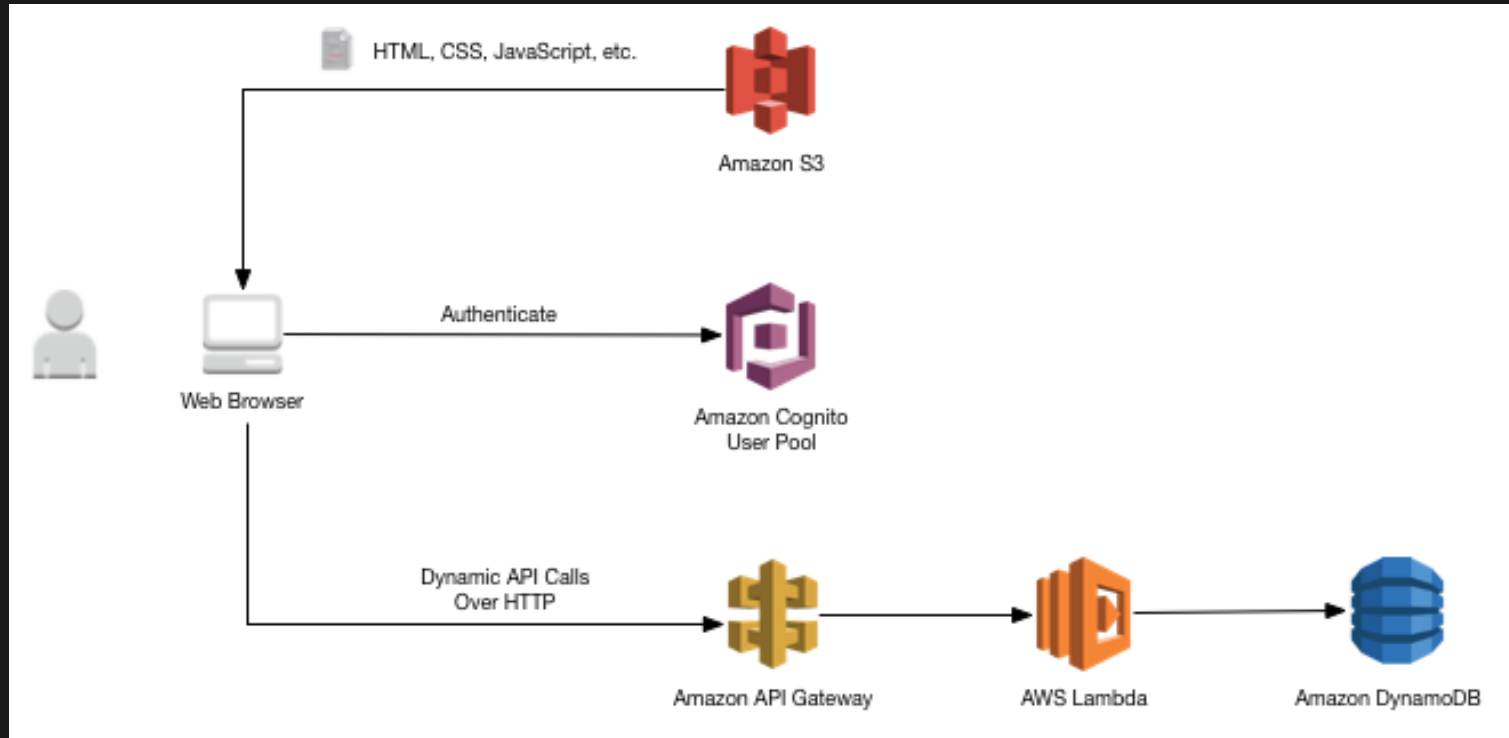
Your Task: Build the Wild Rydes Website

Welcome to Wild Rydes Inc.,
Employee #3!



Scenario: Wild Rydes

The Wild Rydes Serverless Web Application Workshop introduces the basics of building web applications using serverless infrastructure.



Lab 1: Static Website Hosting

OBJECTIVE: Create a bucket in Amazon S3 and configure it for static website hosting. The static HTML, JS, and CSS will be served directly to user browsers from Amazon S3.



Static website hosting

Endpoint : <http://wildrydes-johndoe.s3-website-us-west-2.amazonaws.com>

☒ Use this bucket to host a website [Learn more](#)

Index document [i](#)

Error document [i](#)

Redirection rules (optional) [i](#)

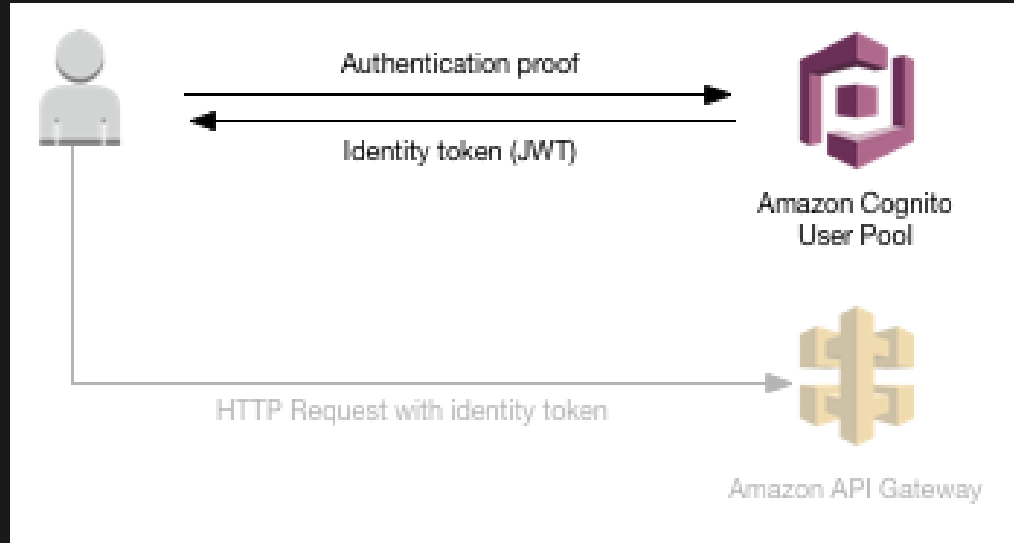
☐ Redirect requests [Learn more](#)

☐ Disable website hosting

[Cancel](#) [Save](#)

Lab 2: User Management

OBJECTIVE: Allow visitors to register as a new user on Wild Rydes, by providing and validating their email address. Amazon Cognito will be used to manage the User Pool for Wild Rydes.



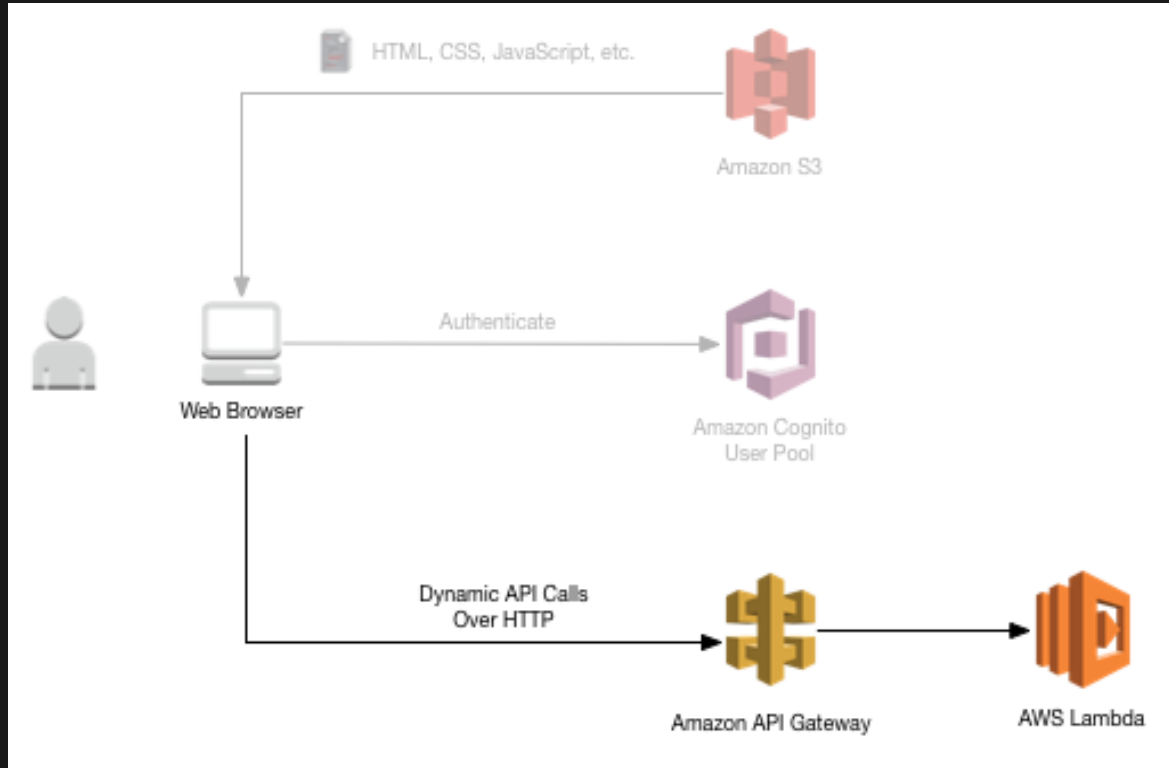
Lab 3: Serverless Service Backend

OBJECTIVE: Create a service backend using AWS Lambda and Amazon DynamoDB to handle requests from your frontend static website content.



Lab 4: Create RESTful API

OBJECTIVE: Use API Gateway to expose the Lambda function you built in the previous module as a RESTful API



WIFI : AWS-Meetup

Password : meetup123

<https://github.com/roryp/wildrydes>